

Getting started with graphics

Andreas Alfons Pieter Schoonees

¹Erasmus School of Economics, Erasmus Universiteit Rotterdam

²Rotterdam School of Management, Erasmus Universiteit Rotterdam

EQI Programming & Visualization for Business Analytics,
September 7, 2018



Course documents and software

- Canvas
- R package [euR](#) contains additional functions and data sets used in this course

Make sure to have the [latest versions](#) of the following software installed:

- [R](http://CRAN.R-project.org): <http://CRAN.R-project.org>
- [RStudio](http://RStudio.com): <http://RStudio.com>



About the lecturer

Andreas Alfons

- PhD in Statistics from Vienna University of Technology
- Assistant professor at the Econometric Institute, ESE
- Research interests:
 - Robust statistics
 - Computer-intensive methods
 - Predictive modeling
- R user and developer for 10+ years
- Austria



Lecture structure

- Materials: Lecture slides with R script; exercises with R script solutions (afterwards)
 - Also, additional material from [R for Data Science](#) book
- Lecture for 20 min, exercises for 20 min, repeat
 - Practice (at home) makes perfect!
- Lecturers can help if you get stuck
 - Do not fall behind
- Learn by helping each other
 - Also on the [discussion board](#) on Canvas
- Consult (1) course material and (2) book before other sources
 - There are plenty of ways to skin a cat



Assignment

Two parts:

- 1 Midterm assignment after week 4
→ Due 7 October 2018
- 2 Final assignment after week 7
→ Due 28 October 2018

→ Similar to the in-class / take home exercises



Content

- 1 About R
- 2 Getting started with R and RStudio
- 3 Packages, data and help
- 4 Built-in R graphics versus package ggplot2
- 5 Basic plots
- 6 Conclusions



About R



About R

- Open source environment for statistical computing
- Free as in “free speech” and “free beer”
- Cross platform: Linux/Unix, Mac OS X, Microsoft Windows
- Fully developed and easy-to-use programming language
- Extensible by community contributed packages
 - Roughly 16000 packages on CRAN, Bioconductor and Github
- Support for businesses by Microsoft:
<https://www.microsoft.com/en-us/cloud-platform/r-server>

→ More information on <http://www.R-project.org/>



R books

- Wickham and Grolemund (2017): [R for Data Science](#)
(read it online for free on <http://r4ds.had.co.nz/>)
- Kabacoff (2015): [R in Action](#)
- Wickham (2009): [ggplot2: Elegant Graphics for Data Analysis](#)
- Specialized topics:
 - Springer's [Use R!](#) series
 - Chapman & Hall/CRC's [The R Series](#)



R resources

- Interactive online learning platform [DataCamp](http://www.DataCamp.com):
<http://www.DataCamp.com>
- Tutorial collection [Quick-R](http://www.statmethods.net): <http://www.statmethods.net>
- Aggregate news and tutorial site [R-bloggers](http://www.R-bloggers.com):
<http://www.R-bloggers.com>
- Scientific articles on R software:
 - [Journal of Statistical Software](http://www.jstatsoft.org): <http://www.jstatsoft.org>
 - [R Journal](http://journal.R-project.org/): <http://journal.R-project.org/>
- R for Matlab users (comparison of R and Matlab commands):
<http://mathesaurus.sourceforge.net/octave-r.html>



Design principles

- Interactive data analysis via **command line interface (CLI)**
 - Enter command and resulting output is printed
- Easy transition from user to developer
 - **Integrated development environment (IDE)**: e.g., RStudio
- **Vectorized arithmetic**
 - Scalars are vectors of length 1
- **Data frames** for storing variables of different types
- **Formula interface** for model specification: $y \sim x_1 + x_2$



What is R used for?

1 Collecting data

- Scrape from the web, import from databases, ...

2 Preparing, exploring and cleaning data

- Data wrangling, exploratory data analysis, [plotting](#)

3 Modelling

- Regression, segmentation, machine learning, custom methods, ...

4 Model evaluation

- Assessing model quality

5 Reporting results

- Writing (dynamic) reports, visualization, creating dashboards, ...

→ And various other tasks



Who Uses R?

- Google
 - Twitter
 - Facebook
 - New York Times
 - John Deere
 - Deloitte
 - Credit Suisse
 - Novartis
 - eBay
 - Ford Motor Company
 - Kickstarter
 - Uber
 - Airbnb
 - Booking.com
 - Bank of America
 - McKinsey & Company
 - FourSquare
- You too?




Who Uses R?

- Google
 - Twitter
 - Facebook
 - New York Times
 - John Deere
 - Deloitte
 - Credit Suisse
 - Novartis
 - eBay
 - Ford Motor Company
 - Kickstarter
 - Uber
 - Airbnb
 - Booking.com
 - Bank of America
 - McKinsey & Company
 - FourSquare
- You too?



Target



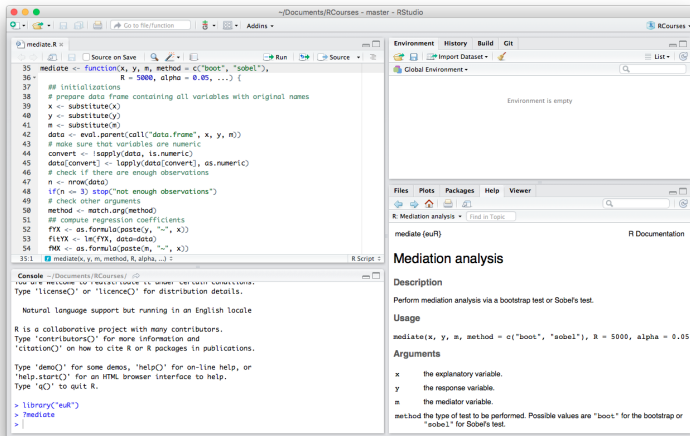
→ After week 2

Getting started with R and RStudio



RStudio

→ Available from <http://RStudio.com>



RStudio: four panels

Top left **Script editor**

Bottom left **R console**

Top right **Two tabs**

- **Environment**: list of objects used in the session
- **History**: allows to re-run previous commands

Bottom right **Five tabs**

- **Files**: browse through files on the computer
- **Plots**: graphics are displayed here
- **Packages**: list of installed packages
- **Help**: R help files are displayed here
- **Viewer**: local web content created in the session



How do I use R with RStudio?

Type commands into an R script and execute from there:

1 Create objects

- Typically contain data, or statistics derived from data
- For example, data loaded from a spreadsheet
- You have to pick (informative) names

2 Manipulate these objects to create new ones

- E.g., transform data, fit a model or create a plot
- All actions are performed by functions

3 Interpret and report your results



Example session 1

Open RStudio, and open `my_first_script.R` in the script editor

- Available on Blackboard
- Execute the line in which your cursor is with *ctrl+enter*
- You can also type the command directly in the console, but it is better to store your commands in a script (reproducibility)



Using R

To use R (efficiently) you need to understand:

- 1 Different types of objects and how they behave
- 2 Various functions and what they do

→ This takes time, but it is worth it!



Functions & objects

- All actions are performed by functions:
 - We used `c()`, `mean()` and `'<-'` (an operator function)
 - Takes an object (or objects), and returns a transformed result
- Functions are pieces of code that perform specific actions
 - For example, `mean()` calculates the mean
 - Function usage is described in its help file (for example, `?mean`)
- Functions accept objects as arguments
 - Data to compute on and other options that control output
 - For example, `age` is a vector object containing ages




Workspace

Created objects are available in your workspace, also known as the Global Environment:



You can save the workspace when you exit R:



→ Better not to: store the code that created the objects



Some details

- R is case sensitive
- The + prompt means R is waiting for you to complete the command
- Press Esc to cancel the command being evaluated
- Use the Tab key for code completion
- Remember to close your parentheses ()
- Press the up and down arrows to browse the command history in the console



Script editor

- Idea: perform an analysis and save the commands in a script
- Execute current line or selection on R console with keyboard shortcut *ctrl+enter*

Advantages and disadvantages:

- + Reproducibility
- + Easy to perform the same analysis with different data
 - Just change the input data
 - Particularly useful, e.g., for periodic reports
- Steeper learning curve than GUI



Script files

Create new script file: *File* → *New file* → *R Script*

Save script file:

- Keyboard shortcut: *ctrl+S* (Windows/Linux), *cmd+S* (Mac)
- *File* → *Save As...* and enter the file name in the dialog

→ Use file extension *.R*

Open existing script file:

- *File* → *Open File...* and select the script file in the dialog
- Click the *Files* tab in the lower right panel, navigate to the script file and click on it



Packages, data and help

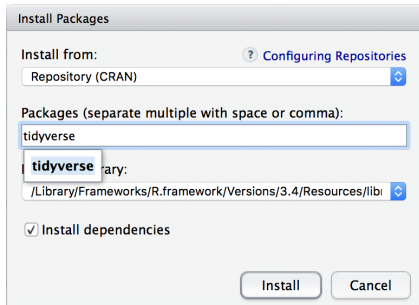


Install packages from CRAN

Tools → *Install Packages...*

In the dialog:

- In the *Install from:* box, select *Repository*
- Type the names of the packages into the text box (suggestions are shown as you type)
- Make sure that *Install dependencies* is checked
- Click *Install*




Install local packages

Tools → *Install Packages...*

In the dialog:

- In the *Install from:* box, select *Package Archive File*
- Click *Browse* and select the package file
- Click *Install*



Install packages: command line

From CRAN:

```
R> install.packages("tidyverse")
```

Local packages:

Windows:

```
R> install.packages("euR_0.1.2.zip", repos = NULL)
```

Mac:

```
R> install.packages("euR_0.1.2.tgz", repos = NULL)
```

Linux/Unix:

```
R> install.packages("euR_0.1.2.tar.gz", repos = NULL)
```

→ If the package is not in the working directory, the **full path** needs to be specified



Load installed packages

In RStudio:

- 1 In the lower right panel, click the *Packages* tab
- 2 Check the box next to the packages to be loaded

On the command line:

```
R> library("euR")
```

- Install a package once on a computer
- Load it in every new session



R help

- R comes with **built-in help facility** to get more information on functionality
- Help topic is usually the name of a function, data set or package

Contributed packages:

- Help files are **required** for packages on CRAN
- Package needs to be loaded to view its help files
- Overview of all help files within a package is available



View R help files

In RStudio:

- 1 In the lower right panel, click the *Help* tab
- 2 Type the topic into the text box on the right (suggestions are shown as you type)



View R help files from the command line

Help for the `help()` function is available:

```
R> ?help  
R> help("help")
```

List all help topics within a package:

```
R> help(package = "euR")
```

Run examples from a help file:

```
R> example("mean")
```



Load data sets from packages

From a package that is **already loaded**:

```
R> data("Prestige")
```

From a package that is **installed, but not loaded**:

```
R> data("Prestige", package = "euR")
```

List all **data sets** of an installed package:

```
R> data(package = "euR")
```



View loaded data set objects

In RStudio:

- 1 In the top right panel, click the *Environment* tab
- 2 Click on the data set object name

On the command line:

- Type the name of the data set to print it
- Use the `View()` function to open RStudio's viewer



View data sets: command line

→ **Example:** prestige of occupations in Canada

```
R> Prestige
```

	education	logincome	women	prestige	type
gov.administrators	13.11	13.59234	11.16	68.8	prof
general.managers	12.26	14.65949	4.02	69.1	prof
accountants	12.77	13.17851	15.70	63.4	prof
----- (output removed from slides) -----					
longshoremen	8.37	12.21462	0.00	26.1	bc
typesetters	10.00	12.65777	13.58	42.2	bc
bookbinders	8.55	11.82058	70.87	35.2	bc

→ **Too much output** even for moderately sized data sets



View first rows of data: head()

→ Get overview of what the data look like

```
R> head(Prestige)
```

	education	logincome	women	prestige	type
gov.administrators	13.11	13.59234	11.16	68.8	prof
general.managers	12.26	14.65949	4.02	69.1	prof
accountants	12.77	13.17851	15.70	63.4	prof
purchasing.officers	11.42	13.11390	9.11	56.8	prof
chemists	14.62	13.03669	11.68	73.5	prof
physicists	15.64	13.42915	5.13	77.6	prof

View last rows of data: tail()

```
R> tail(Prestige)
```

	education	logincome	women	prestige	type
train.engineers	8.49	13.11065	0.00	48.9	bc
bus.drivers	7.58	12.44139	9.47	35.9	bc
taxi.drivers	7.93	12.04439	3.59	25.1	bc
longshoremen	8.37	12.21462	0.00	26.1	bc
typesetters	10.00	12.65777	13.58	42.2	bc
bookbinders	8.55	11.82058	70.87	35.2	bc

Summarize data: summary()

→ Get overview of the **marginal distributions** of the variables

```
R> summary(Prestige)
```

education	logincome	women
Min. : 6.380	Min. : 9.255	Min. : 0.000
1st Qu.: 8.445	1st Qu.:12.003	1st Qu.: 3.592
Median :10.540	Median :12.534	Median :13.600
Mean :10.738	Mean :12.494	Mean :28.979
3rd Qu.:12.648	3rd Qu.:12.999	3rd Qu.:52.203
Max. :15.970	Max. :14.659	Max. :97.510

prestige	type
Min. :14.80	bc :44
1st Qu.:35.23	wc :23
Median :43.60	prof:31
Mean :46.83	NA's: 4
3rd Qu.:59.27	
Max. :87.20	

Loading .RData files

- RStudio: Click on the data file in the *Files* tab
- Command line: Use `load("path/to/file")`



Working directory

R follows a **one directory per project** philosophy

→ Location where R starts looking for files on the file system

In RStudio:

1 *Session* → *Set Working Directory* → *Choose Directory...*

2 In the dialog, select the desired working directory

You can automate this by using **RStudio projects**:

→ Opening the RStudio project restores the working directory



Working directory

On the command line:

```
R> setwd("../foo/bar")
```

- Better to use **relative path** (with respect to working directory) than absolute path (e.g., C:/Users/andreas/foo/bar)
- Go back to parent directory with ..
- Always use / as path separator (instead of \ on Windows)



Exercises

Download and open the “Graphics-Exercises.pdf” file from Blackboard, and do Exercises 1.1 and 1.2



Built-in R graphics versus package ggplot2



The usual suspect

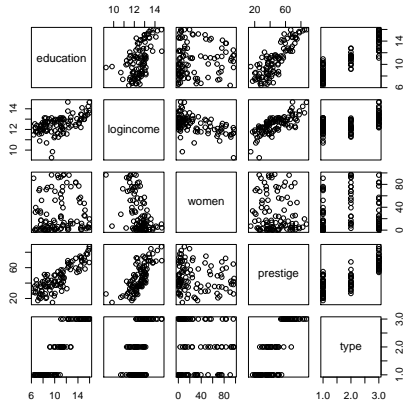
Function `plot()`:

- Scatterplot matrix for data frame
- Works with many other objects, e.g., density estimates, linear models
- Whatever analysis you do, always check if you can `plot()` the result



Scatterplot matrix

```
R> plot(Prestige)
```



Built-in R graphics

- + Allow the user to create quick **plots for exploring the data**
 - + Easy to **add elements** to an existing plot
 - + Fine tuning to produce **high-quality graphics** for publications
 - Designed in the 1970s/80s
 - Sometimes inconsistencies in usage or behavior
 - Customization via cryptic graphical parameters (see ?par)
- Murrell (2011): **R Graphics**



The grammar of graphics

- Designed with recent research on data visualization and human perception in mind
 - Focused on coherence between geometry of the data and geometry of the plot
- The visual representation should fit the data
- Always need to explicitly specify what variables to use and how to plot them
- Implemented in `package ggplot2`
- Wickham (2009): `ggplot2: Elegant Graphics for Data Analysis`



Package ggplot2

- + Coherent approach to graphics
- + Highly flexible and customizable via options and layers
- + Pretty plots (subjective)
- Steeper learning curve than built-in R graphics
- Often not straightforward to add elements to the plot
- Slow even for moderately sized data sets



Basic usage of ggplot2

Add together two basic elements:

- 1 Scaffolding defined by `ggplot()`
 - Selects the data set
 - Defines the variables to be used (the aesthetic mapping):
function `aes()`
- 2 Any number of visual representations of the data, known as `geoms`
 - Define the visual representation (the geometric objects):
function family `geom_xxx()`
 - Different elements are added to the plot using the `+` operator

```
R> library("ggplot2")
```




Basic plots




Scatterplot: Scaffolding

```
R> ggplot(Prestige, aes(x = prestige, y = logincome))
```




Scatterplot: Scaffolding + points

```
R> ggplot(Prestige, aes(x = prestige, y = logincome)) + geom_point()
```



Histogram

```
R> ggplot(Prestige, aes(x = prestige)) + geom_histogram()
```




Histogram: number of bins

- For histograms, it is always a good idea to play with the number of bins
- Number of bins can be specified with argument `bins`
- Bin width can be specified with argument `binwidth`




Histogram: number of bins

```
R> ggplot(Prestige, aes(x = prestige)) + geom_histogram(bins = 15)
```



Density plot

```
R> ggplot(Prestige, aes(x = prestige)) + geom_density()
```




Density plot: kernel and bandwidth

- Density estimate depends on the kernel and smoothing bandwidth
 - Default Gaussian kernel is symmetric and therefore not optimal for asymmetric distributions
- Still useful to get an insight on the shape of the distribution, but be aware of those issues



Quantile-quantile plot

```
R> ggplot(Prestige, aes(sample = prestige)) + geom_qq()
```




Quantile-quantile plot: straight line?

- Plot sample quantiles against theoretical quantiles
- If the distributional assumption holds, the points form almost a straight line
- By default the normal distribution is used
- Distribution can be specified with argument `distribution`



Boxplot

```
R> ggplot(Prestige, aes(x = "", y = prestige)) + geom_boxplot()
```



Boxplot statistics

Upper whisker Largest point still within $1.5 \cdot IQR$ of the upper quartile

Top of box Upper quartile (i.e., 75% quantile)

Middle line Median (i.e., 50% quantile)

Bottom of box Lower quartile (i.e., 25% quantile)

Lower whisker Smallest point still within $1.5IQR$ of the lower quartile

IQR Interquartile range (i.e., difference between upper and lower quartile)


→ No assumption about statistical distribution

→ But: definition of whiskers assumes some degree of symmetry




Conditional boxplot

```
R> ggplot(Prestige, aes(x = type, y = prestige)) + geom_boxplot()
```




Conditional boxplot: significant differences?

```
R> ggplot(Prestige, aes(x = type, y = prestige)) +  
+   geom_boxplot(notch = TRUE)
```



Barplot

```
R> ggplot(Prestige, aes(x = type)) + geom_bar()
```



Time series plot

→ Simply use `geom_line()` instead of `geom_point()` to draw connected line instead of scattered points


→ **Example:** Eredivisie points of Feyenoord

```
R> data("Feyenoord")
```



Time series plot

```
R> ggplot(Feyenoord, aes(x = Year, y = Points)) + geom_line()
```



Some geoms

For a complete list of geoms, [click here](#). Important ones include:

<code>geom_point()</code>	Points
<code>geom_line()</code>	Lines / time series
<code>geom_hline()</code>	Horizontal lines
<code>geom_vline()</code>	Vertical lines
<code>geom_bar()</code>	Bars
<code>geom_boxplot()</code>	Box and whiskers plot
<code>geom_density()</code>	Density estimate
<code>geom_smooth()</code>	Fitted regression line
<code>geom_text()</code>	Text
<code>geom_label()</code>	Text within rectangle
<code>geom_tile()</code>	Rectangles for heat maps

→ Use appropriate geoms!



Exercises

Load the Patents data from the `patents.RData` file, and do Exercise 2.1.



Conclusions



Conclusions

- R/RStudio are **open source** solutions for statistical analysis with a **large community** of users
- Command line interface has **steeper learning curve**, but offers **greater flexibility**
- Basic function **ggplot()** to initialize the plot
- Function **aes()** to define the variables to be used
- Function family **geom_XXX()** to define the visual representation
- Use scripts for **reproducibility** of the plots



References

- R.I. Kabacoff. **R in Action**. Manning, 2nd edition, 2015.
- P. Murrell. **R Graphics**. Chapman & Hall/CRC, 2nd edition, 2011.
- H. Wickham. **ggplot2: Elegant Graphics for Data Analysis**. Springer-Verlag, 2009.
- H. Wickham and G. Grolemund. **R for Data Science**. O'Reilly, 2017.

